

# CS747 – Assignment 1 Report

## Bandit algorithms

Mahesh Abnave (203059010)

### 1. Algorithms

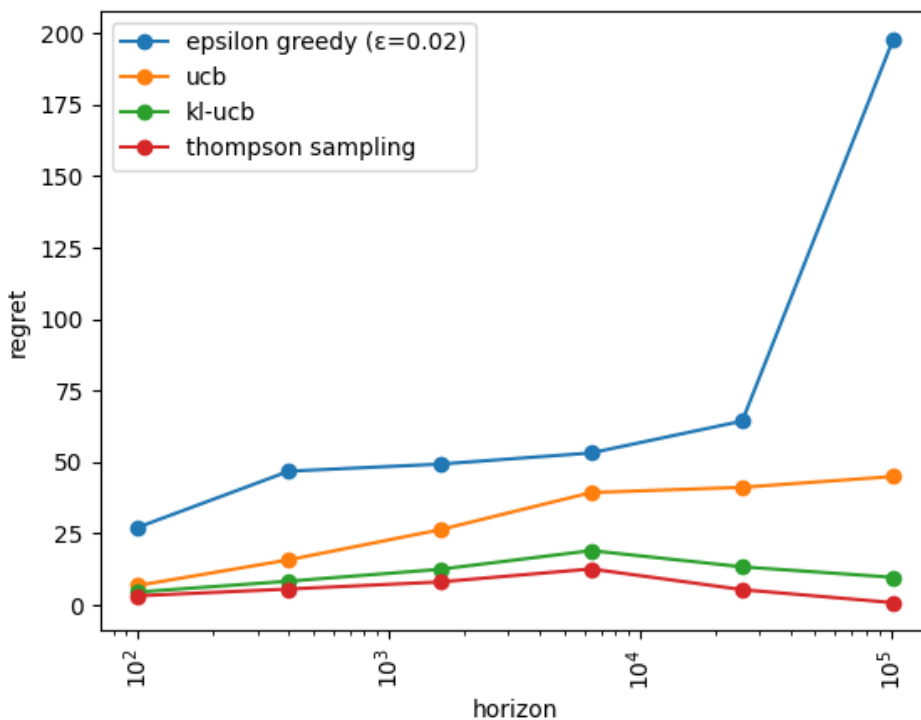
- Implemented following three algorithms:
  1. Epsilon greedy
  2. UCB
  3. KL-UCB
  4. Thompson sampling
  5. Thompson sampling with hint

### 2. Code

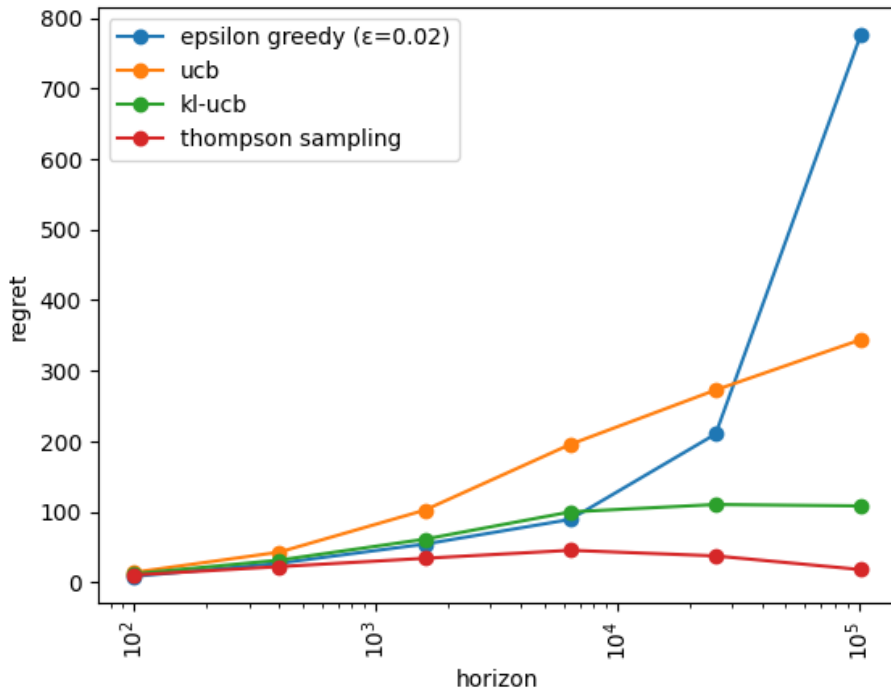
- There are py files corresponding to each algorithm.
- Bandit\_instace.py emulates pulling arm once initialized with weights.
- For number of horizons equal to number of arms, round robin approach is followed to pull each arm once.

### 3. Task 1 – Epsilon greedy vs UCB vs KL-UCB vs Thompson sampling

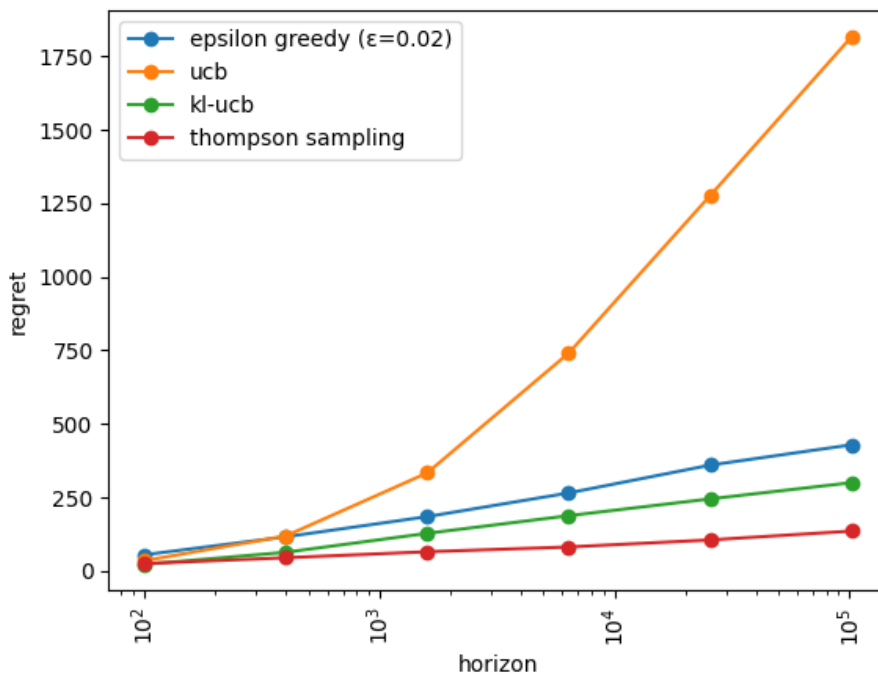
#### 3.1. Plots



**Fig. Evaluating regrets of vairous algorithms for instance 1**



**Fig. Evaluating regrets of vairous algorithms for instance 2**



**Fig. Evaluating regrets of vairous algorithms for instance 3**

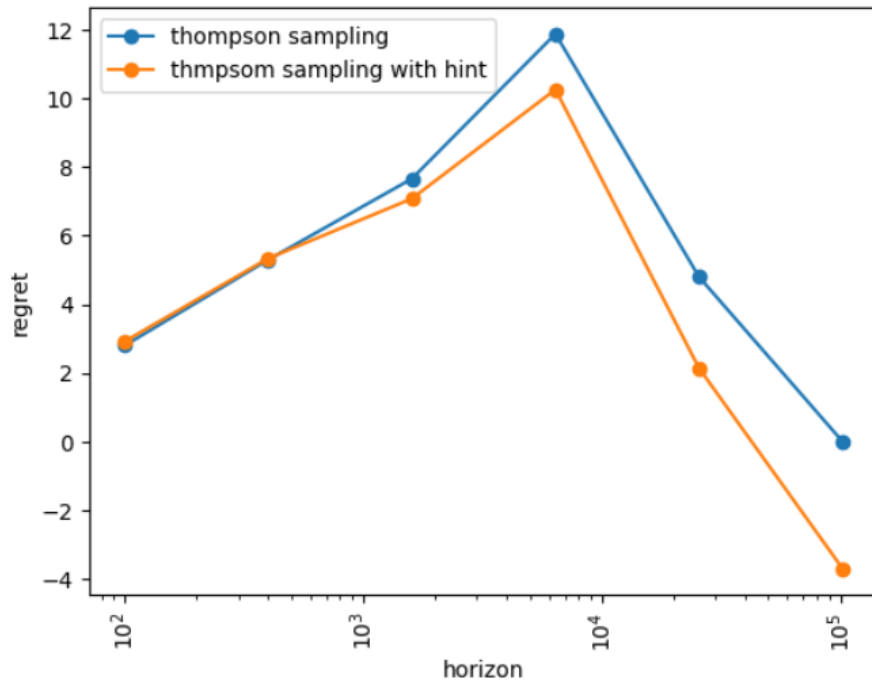
### 3.2. Observations

- All algorithms gives sublinear regret.
- Thompson sampling performs best in all bandit instances.
- KL-UCB behaves second best in all bandit instances.

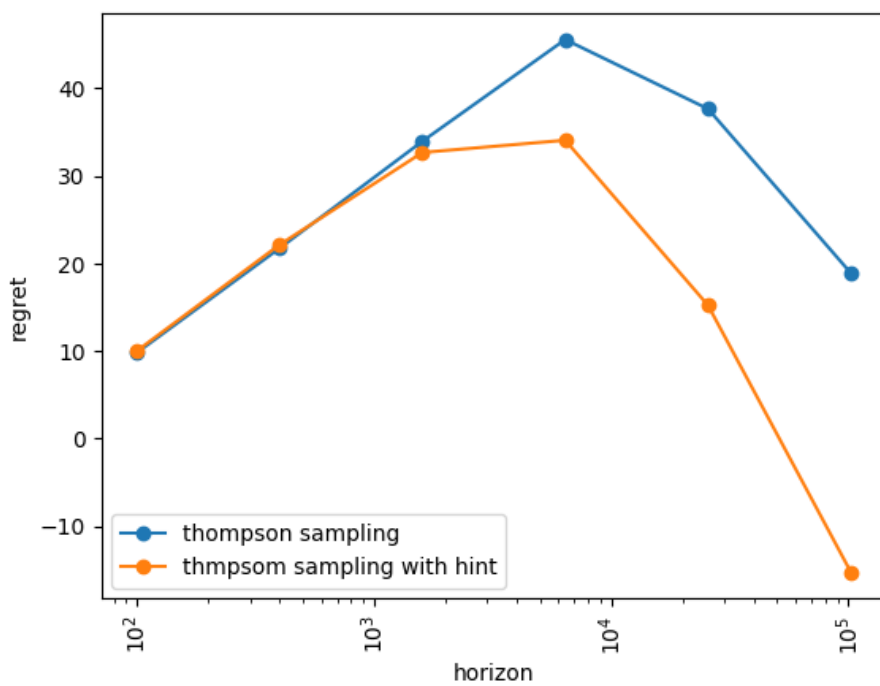
- UCB seems to take time to converge to true mean as number of arms increases. (As for two arm bandits, it performs better than epsilon greedy. For five arm bandit, it initially performs poorer than epsilon greedy, but eventually ends up performing better than it.)

## 4. Task 2 - Thompson sampling vs Thompson sampling with hint

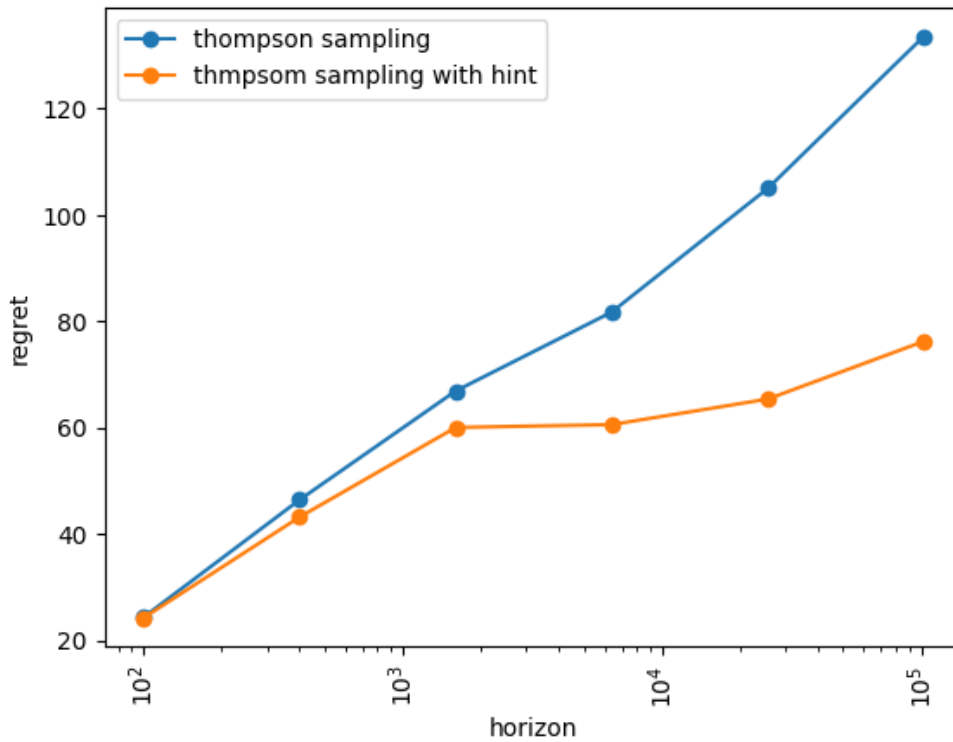
### 4.1. Plots



**Fig. Evaluating regrets of various algorithms for instance 1**



**Fig. Evaluating regrets of various algorithms for instance 2**



**Fig. Evaluating regrets of various algorithms for instance 3**

#### 4.2. Idea of Thompson Sampling with hint algorithm

The idea is simple. For initial iteration, we startup with normal Thompson sampling algorithm. And once we find optimal arm, we keep exploiting it. For finding optimal arm, we perform additional tasks in each iteration. Say for first iteration, we sample according to normal Thompson sampling algorithm. Then we check which arm has empirical mean (EM) closest to max true mean (MTM). We “note down” index of this arm. In addition, we increment counter which tells how many “last consecutive” times, an arm is found to have EM closest to MTM. So if in second iteration, we find that the same arm has EM closest to MTM, we increment this counter. Now, say counter is 100 and we find new arm to have EM closest to MTM. Then we note down index of this new possibly-optimal arm and reset counter to 1. We assume arm is optimal once its found to have EM closest to MTM for (threshold) 1000 consecutive iterations. Then onwards we simply exploit this possibly-optimal arm (that is always pull this arm instead of one as per naïve Thompson sampling). Once arm is found to be optimal, we still keep checking if its EM falls below EM of other arms. And if this happens, we reset the whole process again.

#### 4.3. Pseudocode for Thompson sampling with hint algorithm

(# green text is comment)

1. **Initialise:**  
last\_arm\_with\_em\_closest\_to\_max\_tm = -1  
no\_of\_times\_arm\_has\_em\_closest\_to\_max\_tm = 0 #for last "consecutive" iterations t  
optimal\_arm\_determined = False  
max\_true\_mean = max(true\_means)
2. **For** t = 1 **to** no\_of\_arms: #initially pull each arm once  
    **call** pull\_arm\_and\_update\_values(arm\_to\_pull = t)
3. **For** t > no\_of\_arms:  
    **if** NOT optimal\_arm\_determined:  
        beta\_per\_arm = sample all arms from beta distribution  
        arm\_to\_pull = argmax(beta\_per\_arm)  
    **else:**  
        arm\_to\_pull = last\_arm\_with\_em\_closest\_to\_max\_tm  
    **call** pull\_arm\_and\_update\_values(arm\_to\_pull)
4. **function** pull\_arm\_and\_update\_values(arm\_to\_pull)  
    number\_of\_pulls(arm\_to\_pull)++  
    **if** reward == 1:  
        number\_of\_successes[arm\_to\_pull]++  
    **else:**  
        number\_of\_failures[arm\_to\_pull]++  
    empirical\_means[arm\_to\_pull] = #EM formula  
    **if** NOT optimal\_arm\_determined:  
        arm\_with\_em\_closest\_to\_max\_tm = #arm which has its EM closest to max TM  
        **if** last\_arm\_with\_em\_closest\_to\_max\_tm == arm\_with\_em\_closest\_to\_max\_tm:  
            # arm with its EM closest to max TM is same as that in last iteration t  
            no\_of\_times\_arm\_has\_em\_closest\_to\_max\_tm++  
        **else:**  
            # arm with its EM closest to max TM is different than that in last iteration t  
            last\_arm\_with\_em\_closest\_to\_max\_tm = arm\_with\_em\_closest\_to\_max\_tm  
            no\_of\_times\_arm\_has\_em\_closest\_to\_max\_tm = 1  
        **if** no\_of\_times\_arm\_has\_em\_closest\_to\_max\_tm == 1000:  
            # if arm is found to have EM closest to max TM for last 1000 "consecutive"  
            # iterations, assume it's the optimal arm  
            optimal\_arm\_determined = True  
    **else:** # that is if optimal\_arm\_determined = True  
        # check if the arm which we assumed to be optimal is really optimal  
        arm\_with\_em\_closest\_to\_max\_tm = #arm which has its EM closest to max TM  
        **if** last\_arm\_with\_em\_closest\_to\_max\_tm == arm\_with\_em\_closest\_to\_max\_tm:  
            # arm which we assumed to be optimal turned out to be sub-optimal, reset  
            # everything  
            optimal\_arm\_determined = False  
            last\_arm\_with\_em\_closest\_to\_max\_tm = arm\_with\_em\_closest\_to\_max\_tm  
            no\_of\_times\_arm\_has\_em\_closest\_to\_max\_tm = 1

#### 4.4. Observations


- The Thompson sampling with hint algorithm formulated seem to perform better on all instances.
- Due to threshold 1000, it performs similar to (or “slightly” better than) naïve Thompson sampling. But after  $t=1000$ , it starts performing significantly better than naïve Thompson sampling, because of exploiting of optimal arm.

## 5. Task 3 - Experiments on epsilon-greedy

### 5.1. Observations

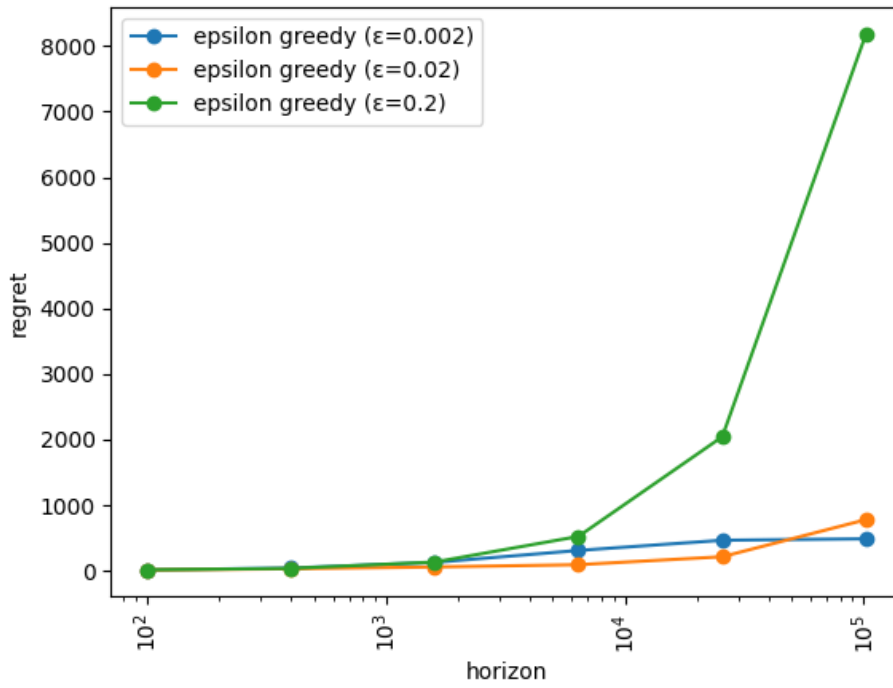
For  $\epsilon = 0.002$ ,  $\epsilon = 0.02$  and  $\epsilon = 0.2$ , it was found that epsilon greedy algorithm with  $\epsilon = 0.02$  performs better than the other two with significant margin except for instance 2. For instance 2, it performs “slightly” poorer than  $\epsilon = 0.002$ , that too for higher horizon. For instance 2, for not so bigger horizon ( $< 10^5$ ),  $\epsilon = 0.02$  still performs better than the other two.

### 5.2. Plots




**Fig. Evaluating regrets of various algorithms for instance 1**





**Fig. Evaluating regrets of various algorithms for instance 2**



**Fig. Evaluating regrets of various algorithms for instance 3**